

Introduction to Scikit-Learn: Machine Learning with Python

Regression

郭耀仁

Regression

About regression

One of the simplest regression problems is fitting a line to data, which we've seen previously.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression


train_url = "https://storage.googleapis.com/kaggle_datasets/House-Prices-Advanced-Regression-Techniques/train.csv"
train_df = pd.read_csv(train_url)
X_train = train_df["GrLivArea"].values.reshape(-1, 1)
y_train = train_df["SalePrice"].values.reshape(-1, 1)
reg = LinearRegression()
reg.fit(X_train, y_train)
```

```
Out[1]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [2]: xfit = np.linspace(X_train.min() - 10, X_train.max() + 10, 100).reshape(-1, 1)
yfit = reg.predict(xfit)
plt.scatter(train_df["GrLivArea"], train_df["SalePrice"], label='train', s=3, color="#4286f4")
plt.plot(xfit, yfit, color="#f4a041", linewidth=2, label='thetas')
plt.legend()
```

```
Out[2]: <matplotlib.legend.Legend at 0x1a0b5bc470>
```

```
In [3]: plt.show()
```



How to deal with polynomial features

Using `PolynomialFeatures()` function.

```
In [4]: from sklearn.preprocessing import PolynomialFeatures
```

```
poly = PolynomialFeatures(5)  
X_train_poly = poly.fit_transform(X_train)  
print(X_train_poly)
```

```
[[ 1.00000000e+00  1.71000000e+03  2.92410000e+06  5.00021100e+09  
   8.55036081e+12  1.46211170e+16]  
 [ 1.00000000e+00  1.26200000e+03  1.59264400e+06  2.00991673e+09  
   2.53651491e+12  3.20108182e+15]  
 [ 1.00000000e+00  1.78600000e+03  3.18979600e+06  5.69697566e+09  
   1.01747985e+13  1.81721902e+16]  
 ...,  
 [ 1.00000000e+00  2.34000000e+03  5.47560000e+06  1.28129040e+10  
   2.99821954e+13  7.01583371e+16]  
 [ 1.00000000e+00  1.07800000e+03  1.16208400e+06  1.25272655e+09  
   1.35043922e+12  1.45577348e+15]  
 [ 1.00000000e+00  1.25600000e+03  1.57753600e+06  1.98138522e+09  
   2.48861983e+12  3.12570651e+15]]
```




```
In [5]: reg = LinearRegression()  
reg.fit(X_train_poly, y_train)
```

```
Out[5]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [6]: xfit = np.linspace(X_train.min() - 10, X_train.max() + 10, 100).reshape(-1, 1)
xfit_poly = poly.fit_transform(xfit)
yfit = reg.predict(xfit_poly)
plt.scatter(train_df["GrLivArea"], train_df["SalePrice"], label='train', s=3, color="#4286f4")
plt.plot(xfit, yfit, color="#f4a041", linewidth=2, label='thetas')
plt.legend()
```

```
Out[6]: <matplotlib.legend.Legend at 0x1a18177e48>
```

```
In [7]: plt.show()
```



Or using more sophisticated models

In [8]: `from sklearn.ensemble import RandomForestRegressor`

```
rf_reg = RandomForestRegressor()  
rf_reg.fit(X_train, y_train.ravel())
```

Out[8]: `RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
oob_score=False, random_state=None, verbose=0, warm_start=False)`

```
In [9]: xfit = np.linspace(X_train.min() - 10, X_train.max() + 10, 100).reshape(-1, 1)
yfit = rf_reg.predict(xfit)
plt.scatter(train_df["GrLivArea"], train_df["SalePrice"], label='train', s=3, color="#4286f4")
plt.plot(xfit, yfit, color="#f4a041", linewidth=2, label='thetas')
plt.legend()
```

```
Out[9]: <matplotlib.legend.Legend at 0x1a184f2780>
```

```
In [10]: plt.show()
```

